

# HANDLING ERRORS AND DEBUGGING

# LEARNING OBJECTIVES

- Learn about how SYCL handles errors
- Learn about the difference between synchronous and asynchronous exceptions
- Learn how to handle exceptions and retrieve further information
- Learn about the host device and how to use it

## SYCL EXCEPTIONS

- In SYCL errors are handled by throwing exceptions.
- It is crucial that these errors are handled, otherwise your application could fail in unpredictable ways.
- In SYCL there are two kinds of error:
  - Synchronous errors (thrown in user thread) .
  - Asynchronous errors (thrown by the SYCL scheduler).

# HANDLING ERRORS

```
int main() {  
    queue q();  
  
    /* Synchronous code */  
  
    q.submit([&](handler &cgh) {  
  
        /* Synchronous code */  
  
        cgh.parallel_for<add>(buf0.get_range(), [=](id<1> i) {  
  
            /* Asynchronous code */  
  
        });  
    });  
}
```

- Kernels run asynchronously on the device, and will throw asynchronous errors.
- Everything else runs synchronously on the host, and will throw synchronous errors.

# SYCL EXCEPTIONS

Synchronous  
exceptions

SYCL interface

Asynchronous  
exceptions

SYCL Runtime

(optional)  
CPU device

Kernel  
loader

Runtime  
Scheduler

Data dependency  
tracker

Backend interface (e.g. OpenCL API)

# HANDLING ERRORS

```
class add;

int main() {
    queue q();

    /* Synchronous code */

    q.submit([&](handler &cgh) {
        /* Synchronous code */

        cgh.single_task<add>([=](id<1> i) {
            /* Asynchronous code */
        });
    }).wait();
}
```

- Code on the device runs asynchronously
- If errors are not handled, the application can fail:
  - SYCL 1.2.1 application will fail silently.
  - SYCL 2020 provides a default async handler that will call `std::terminate` when an asynchronous error is thrown.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try {
        queue gpuQueue(gpu_selector{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.parallel_for<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

    } catch (...) { /* handle errors */ }
}
```

- Synchronous errors are typically thrown by SYCL API functions.
- In order to handle all SYCL errors you must wrap everything in a try-catch block.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue gpuQueue(gpu_selector{}, async_handler{});
        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.parallel_for<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

        gpuQueue.throw_asynchronous();
    } catch (...) { /* handle errors */
    }
}
```

- Asynchronous errors that may have occurred will be thrown after a command group has been submitted to a queue.
  - To handle these errors you must provide an async handler when constructing the queue object.

u must also call the `throw_asynchronous` or `wait_and_throw`



```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.parallel_for<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

        gpuQueue.throw_asynchronous();
    } catch (...) { /* handle errors */ }
```

- The async handler is a C++ lambda or function object that takes as a parameter an `exception_list`
- The `exception_list` class is a wrapper around a list of `exception_ptr`s which can be iterated over

`exception_ptr`s can be rethrown by passing them to

```
int main() {  
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };  
    try {  
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {  
            for (auto e : eL) { std::rethrow_exception(e); }  
        });  
  
        ...  
  
        gpuQueue.throw_asynchronous();  
    } catch (const std::exception& e) {  
        std::cout << "Exception caught: " << e.what()  
        << std::endl;  
    }  
}
```

- Once rethrown and caught, a SYCL exception can provide information about the error
- The `what` member function will return a string with more details

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try {
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        ...

        gpuQueue.throw_asynchronous();
    } catch (const sycl::exception& e) {
        std::cout << "Exception caught: " << e.what();
        std::cout << " With OpenCL error code: "
        << e.get_cl_code() << std::endl;
    }
}
```

- In SYCL 1.2.1, if the exception has an OpenCL error code associated with it this can be retrieved by calling the `get_cl_code` member function
- If there is no OpenCL error code this will return `CL_SUCCESS`
- SYCL 2020 provides the `error_category_for` templated free function that allows checking for the category of the exception depending on the backend used (e.g. `backend::opencl`), and `e.code().value()` will correspond to the backend error code.

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };

    queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
        for (auto e : eL) { std::rethrow_exception(e); }
    });
    context gpuContext = gpuQueue.get_context();

    try {
        ...
        gpuQueue.wait_and_throw();
    } catch (const sycl::exception& e) {
        if (e.has_context()) {
            if (e.get_context() == gpuContext) {
                /* handle error */
            }
        }
    }
}
```

- The `has_context` member function will tell you if there is a SYCL context associated with the error
- If that returns true then the `get_context` member function will return the associated SYCL context object

# EXCEPTION TYPES

- In SYCL 1.2.1 there are a number of different exception types that inherit from `std::exception`
  - E.g. `runtime_error`, `kernel_error`
- SYCL 2020 only has a single `sycl::exception` type which provides different error codes
  - E.g. `errc::runtime`, `errc::kernel`

# DEBUGGING SYCL KERNEL FUNCTIONS

- Every SYCL 1.2.1 implementation is required to provide a host device
  - This device executes native C++ code but is guaranteed to emulate the SYCL execution and memory model
- This means you can debug a SYCL kernel function by switching to the host device and using a standard C++ debugger
  - For example gdb



- SYCL 2020 only guarantees that a device will always be available, and users can query the `host_debuggable` device aspect to check whether they can use the same functionality as the SYCL 1.2.1 host device

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue hostQueue(aspect_selector<aspect::host_debuggable>(), async_handler{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        hostQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.parallel_for<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        });
        hostQueue.wait_and_throw();
    } catch (...) { /* handle errors */ }
}
```

- Any SYCL application can be debugged on the host device by switching the queue for a host queue
- Replacing the device selector for the `aspect_selector` will ensure that the queue submits all work to the device with the requested aspects, in this case a debuggable device

# QUESTIONS

# EXERCISE

Code\_Exercises/Exercise\_4\_Handling\_Errors/source

Add error handling to a SYCL application for both synchronous and asynchronous errors.