

# Performance Portability in Practice

John Pennycook and Jason Sewall, Intel Corporation

Acknowledgements:

Doug Jacobsen, Google


Tom Deakin and Simon McIntosh-Smith, University of Bristol

# What is “Performance Portability”?

“A **measurement** of an application’s **performance efficiency** for a given problem that can be executed correctly on all platforms in a given set.”

$$\Phi(a, p, H) = \begin{cases} \frac{|H|}{\sum_{i \in H} \frac{1}{e_i(a, p)}} & \text{if } i \text{ is supported } \forall i \in H \\ 0 & \text{otherwise} \end{cases}$$

- Enables objective comparisons of implementations, algorithms, etc
- See our publications for more detail:
  - “[Implications of a Metric for Performance Portability](#)”
  - “[Navigating Performance, Portability and Productivity](#)”
  - “Revisiting a Metric for Performance Portability” (to appear in P3HPC 2021)



An “efficiency cascade” plot, highlighting performance trends across platforms of interest. Plot your own using [these scripts](#).



# What About Productivity?

- Per-platform code increases developer effort:

- Initial development
- Debugging
- Maintenance

- We measure this via “Code Divergence”:

## Codebase 1

### shared.cpp:

```
void whereami()
{
#ifdef CPU
    printf("CPU\n");
#else
    printf("GPU\n");
#endif
}
```

$$\frac{8 \text{ lines} - 6 \text{ shared lines}}{8 \text{ lines}} = 0.25$$

## Codebase 2


### cpu.cpp:

```
void whereami()
{
    printf("CPU\n");
}
```

### gpu.cpp:

```
void whereami()
{
    printf("GPU\n");
}
```

$$\frac{8 \text{ lines} - 0 \text{ shared lines}}{8 \text{ lines}} = 1.0$$

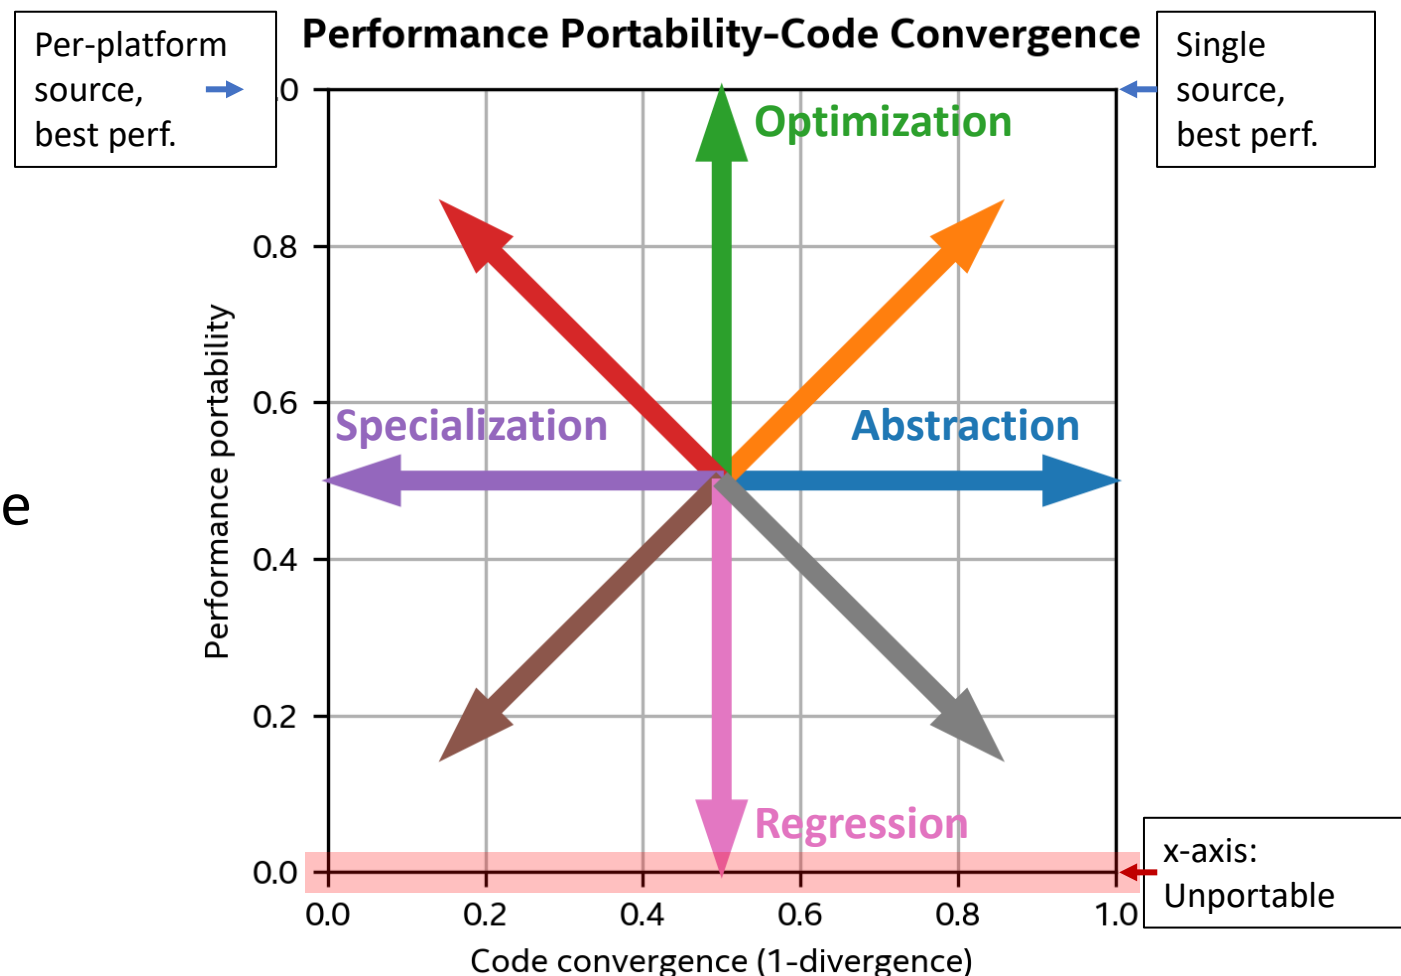


A dendrogram clustering platforms by “code divergence”.  
Plot your own using our [Code Base Investigator](#) tool.



## Setting Goals and Tracking Progress

- Trade-offs abound:
  - **Abstraction**  $\Rightarrow$  Higher Convergence
  - **Specialization**  $\Rightarrow$  Lower Convergence
  - **Optimization**  $\Rightarrow$  Higher Performance
- SYCL provides tools to help us manage these trade-offs
- But we must set our own goals (and manage our expectations!)





## Specialization: Device Queries

*Lower Convergence, Higher Performance*

- Inspect where you're running and dynamically adapt to the capabilities available.

```
sycl::queue q(sycl::default_selector{});
sycl::device d = q.get_device();
// Process input in the largest chunks we can allocate
uint64_t chunk_size = d.get_info<sycl::info::device::max_mem_alloc_size>();
// Prefer double precision if it's supported
if (d.has(sycl::aspect::fp64)) {
    run<double>(chunk_size);
} else {
    run<float>(chunk_size);
}
```

- Many queries, e.g.: cache size, work-group and sub-group sizes, availability of scratchpad memory, atomics capabilities



## Abstraction: Device Selectors

*Higher Convergence, Higher Performance*

- Select the best device for your application

```
struct MySelector : sycl::device_selector {
    // SYCL 2020 allows this to be a lambda
    int operator()(const sycl::device& d) const {
        int score = 0;
        if (not d.has(sycl::aspect::fp64)) {
            score = -1;
        } // Return negative if required features are not supported
        else if (d.has(sycl::aspect::gpu)) {
            score = 1;
        } // Return a higher score for GPU devices
        std::cout << "Scoring: " << d.get_info<sycl::info::device::name>() << " " << score << std::endl;
        return score;
    }
};

int main() {
    // Create a queue associated with the selected device
    sycl::queue q(MySelector{});
    std::cout << "Selected: " << q.get_device().get_info<sycl::info::device::name>() << std::endl;
    ...
}
```



## Abstraction: Subdevices (e.g. NUMA)

*Higher Convergence, Higher Performance*

```
using namespace sycl::info;
constexpr auto by_affinity_domain = partition_property::partition_by_affinity_domain;
auto has_numa_domains = [=](sycl::device &d) {
    auto props = d.get_info<device::partition_properties>();
    auto domains = d.get_info<device::partition_affinity_domains>();
    bool by_affinity = std::find(props.begin(), props.end(), by_affinity_domain) != props.end();
    bool by_numa = std::find(domains.begin(), domains.end(),
                             partition_affinity_domain::numa) != domains.end();
    return by_affinity && by_numa;
};

// Try to split selected device into sub-devices along NUMA boundaries
sycl::device root = sycl::device(sycl::default_selector{});
std::vector<sycl::device> devices;
if (has_numa_domains(root)) {
    devices = root.create_sub_devices<by_affinity_domain>(partition_affinity_domain::numa);
}
else {
    devices = std::vector<sycl::device>{root};
}

std::cout << "Split " << root.get_info<device::name>() << " into " << devices.size()
           << " NUMA domains." << std::endl;
```



## Specialization: Specialization Constants

*Lower Convergence, Higher Performance*

```
q.submit([&](auto &h) {
    h.template set_specialization_constant<BLOCKSIZE>(blocksize);
    h.parallel_for(nd_range<2>({(size_t)N,(size_t)M},{blocksize,blocksize}),
        [=](nd_item<2> it, kernel_handler kh) {
            const auto BS = kh.template get_specialization_constant<BLOCKSIZE>();
            auto id = it.get_group().get_id();
            const int i = id[1]*BS;
            const int j = id[0]*BS;
            if(BS == 4) {
                do_block_static<4>(it, i, j, M, N, K, C, A, B);
            }
            else if(BS == 8) {
                do_block_static<8>(it, i, j, M, N, K, C, A, B);
            }
            else if(BS == 16) {
                do_block_static<16>(it, i, j, M, N, K, C, A, B);
            }
            else {
                do_block_dyn(BS, it, i, j, M, N, K, C, A, B);
            }
        });
}).wait();
```

Specialization constants are runtime values known to be constant during Just-in-Time (JIT) compilation





## Abstraction: Group Algorithms

*Higher Convergence, Higher Performance*

- Syntax and semantics based on ISO C++ algorithms
  - any/all/none\_of, shift\_left/right, reduce, exclusive/inclusive\_scan
  - SYCL specific: permute, select

```
// Joint algorithms divide an explicit range across work-items in a group  
// e.g. reduce elements held in global memory, described by [first, last) iterators  
auto sum = sycl::joint_reduce(it.get_group(), first, last, sycl::plus<>());
```

```
// Algorithms with group prefix/suffix operate on implicit range (the group itself)  
// e.g. reduce x values held in private memory of each work-item  
auto sum = sycl::reduce_over_group(it.get_group(), x, sycl::plus<>());
```

- Available for work-groups and sub-groups



## Specialization: `invoke_simd` (DPC++ Extension)

*Lower Convergence, Higher Performance*

- Exploring ways to provide interoperability between SPMD and SIMD programming

```
int popcount(sycl::ext::oneapi::experimental::simd_mask<bool, 8> mask) {  
    // Utilize SIMD APIs with no SPMD equivalent, e.g., popcount  
    return sycl::ext::oneapi::experimental::popcount(mask);  
}  
...  
q.parallel_for(sycl::nd_range<1>{N},  
               sycl::nd_item<1> it) [[sycl::reqd_sub_group_size(8)]] {  
    // Express algorithm in SPMD where convenient  
    bool cond = (it.get_local_id() % 2);  
    // Pass control to explicitly vectorized SIMD function  
    sycl::sub_group sg = it.get_sub_group();  
    int ct = sycl::ext::oneapi::experimental::invoke_simd(sg, popcount, cond);  
    // Resume SPMD coding  
    ...  
}
```

# Summary and Call to Action

- SYCL exposes many features to help us achieve our desired balance of performance, portability and productivity
- Achieving high levels of performance (portability) may require non-trivial effort; what is “acceptable” effort is subjective
- Help us improve SYCL and implementations:
  - SYCL Specification: <https://github.com/KhronosGroup/SYCL-Docs>
  - DPC++: <https://github.com/intel/llvm>
  - hipSYCL: <https://github.com/illuhad/hipSYCL>
  - Codeplay Proposals: <https://github.com/codeplaysoftware/standards-proposals>